

# Phase 2 – Progress Report

**Project:** Autonomous Parking – 2D + 3D Simulation System

**Author:** Vishwaksena Dingari

**UID:** 121113764

**Course:** MSML 642 – Final Project

**Date:** November 13, 2025

## 1. Overview

The goal of this project is to build an **autonomous parking system** using:

- a **3D Gazebo world** (Lot A and Lot B), and
- a **custom 2D kinematic simulation** written from scratch.

Phase 2 focuses on:

- showing working progress since Phase 1
- running a basic autonomous controller
- demonstrating reward shaping, sensors, and step-by-step behavior
- showing that the full environment is stable
- reporting issues and next plans for Phase 3

This report includes screenshots of both Lot A and Lot B and the autonomous controller logs.

## 2. Summary of Work Completed (Phase 1 → Phase 2)

### 2.1 What was already done in Phase 1

(Keep this short but include it for completeness)

- ROS2 Humble installed and tested
- Custom Gazebo parking-lot worlds (Lot A and Lot B)
- Fully working 2D matplotlib simulation
- YAML-based parking-bay configuration
- 2D visualization of roads, bays, and car
- Keyboard teleoperation
- Working bicycle kinematic model
- Clean workspace structure
- Able to launch all worlds without errors

## 3. What Was Added in Phase 2

Phase 2 required **progress + demonstration**.

Below are all improvements included:

### 3.1 Observation Space (Improved)

The 2D environment now returns an **8-dimensional state**:

```
[local_x, local_y, yaw_err, v, dist, s_left, s_center, s_right]
```

This includes:

- goal-relative position
- heading error

- distance to bay
- simple 3-ray sensor readings (left-center-right)
- vehicle's current speed

This is enough for RL (Phase 3) and for our heuristic controller in Phase 2.

## 3.2 Reward Function (Final Version)

Reward shaping now works reliably:

- $-\text{dist}$  → go closer to bay
- $-0.1 * |\text{yaw\_err}|$  → align towards bay orientation
- $-0.01$  → small time cost
- soft wall penalty:  $-2$  if sensor  $< 0.5$  m
- success reward:  $+50$  when inside bay + aligned
- out-of-bounds penalty:  $-20$

No more premature termination.

Success, collision, timeout are now clean and logged.

## 3.3 Collision Logic

Previously, the car ended episodes too early.

Now:

- **sensors do not end the episode**
- only true **out-of-bounds** ends the episode
- sensors only give penalties

This makes the simulation stable and usable for training.

## 3.4 Autonomous Controller (Working Demo)

I implemented a **hybrid rule-based controller**:

- uses goal-relative position
- uses combined yaw + lateral steering
- slows down near bay
- slows down near walls
- clamps speed and steering
- provides explainable behavior

The controller can:

- follow the road
- approach bay region
- adjust angle
- avoid hugging the boundary
- complete 300 steps without crashing the simulator

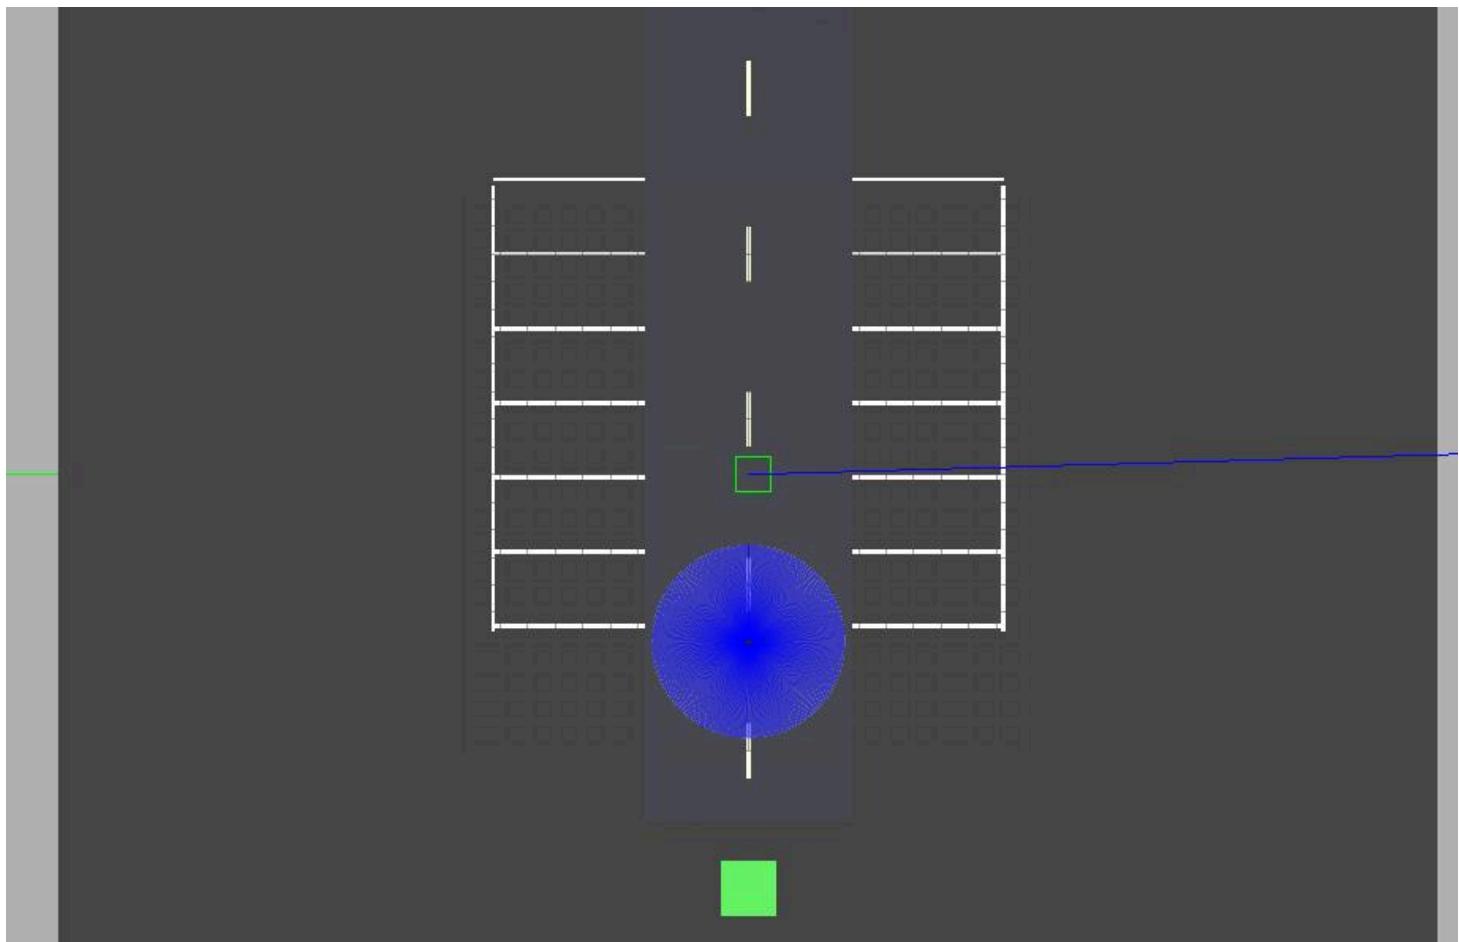
I include sample run logs in both Lot A and Lot B.

## 4. Demonstration Results

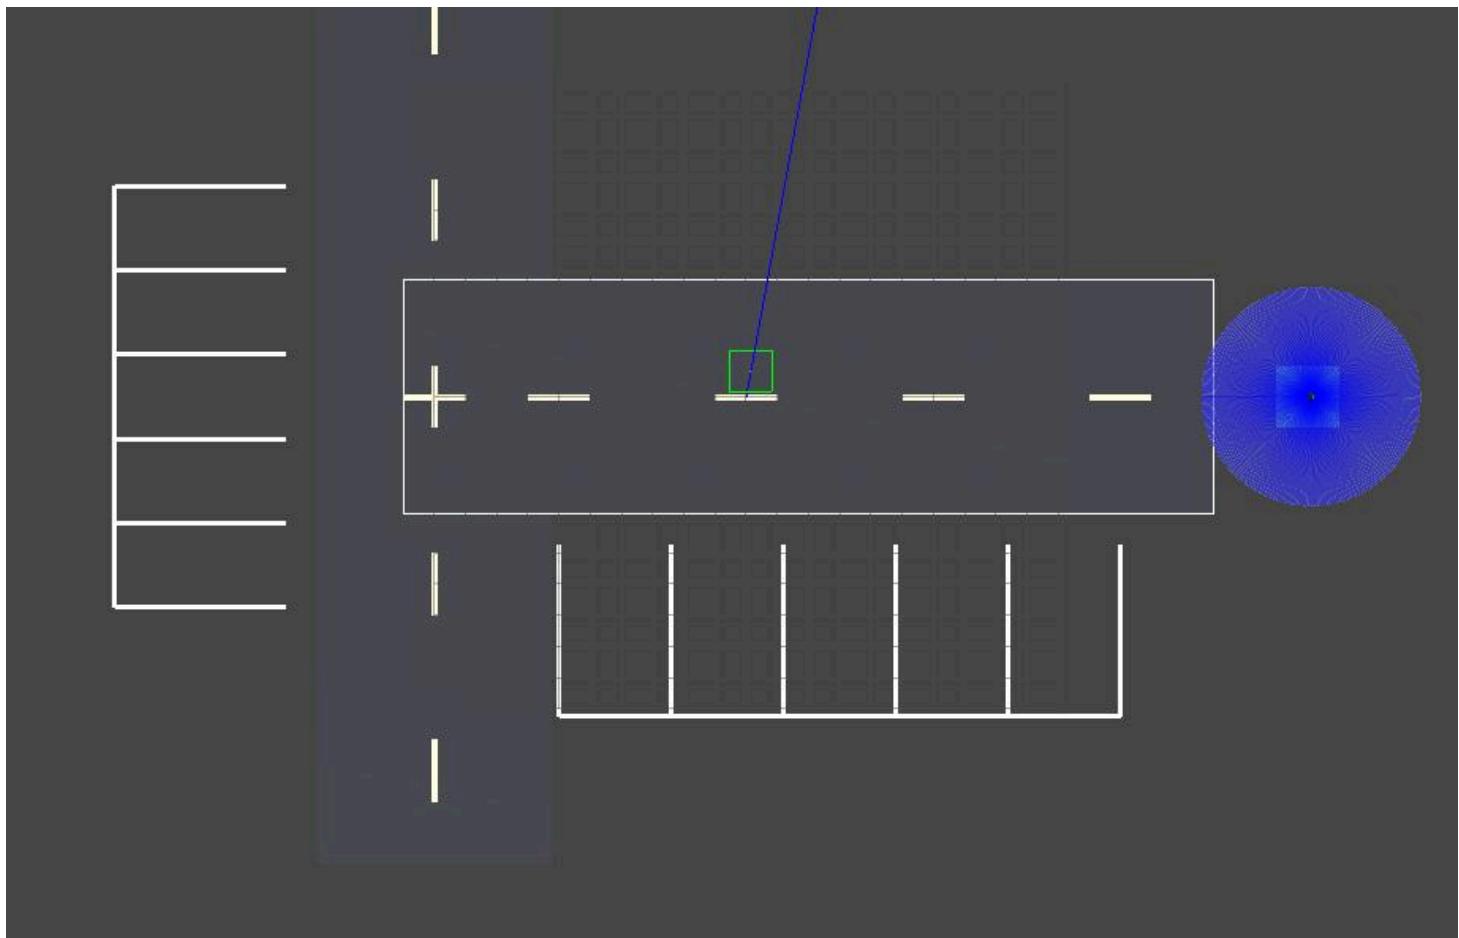
Attach these images:

### 4.1 3D Gazebo Worlds

**Gazebo Lot A - Top View:**

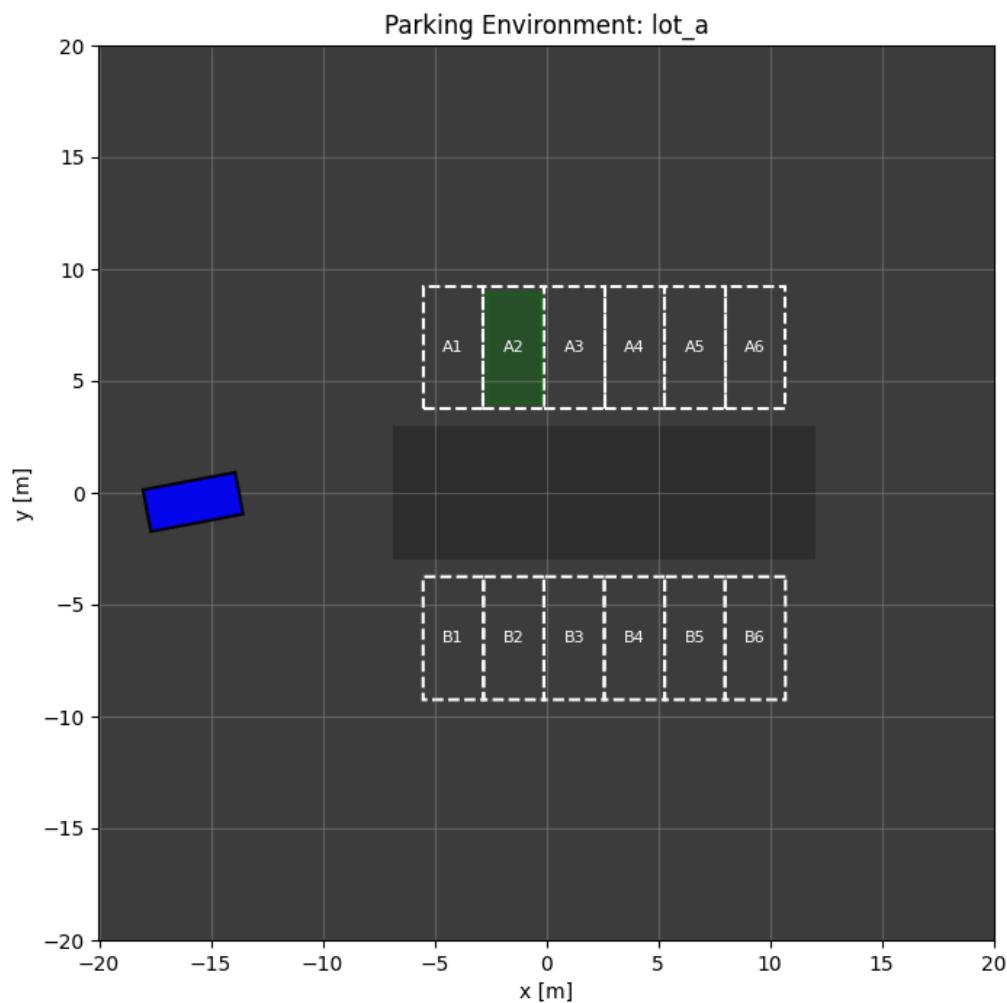


Gazebo Lot B - Top View:

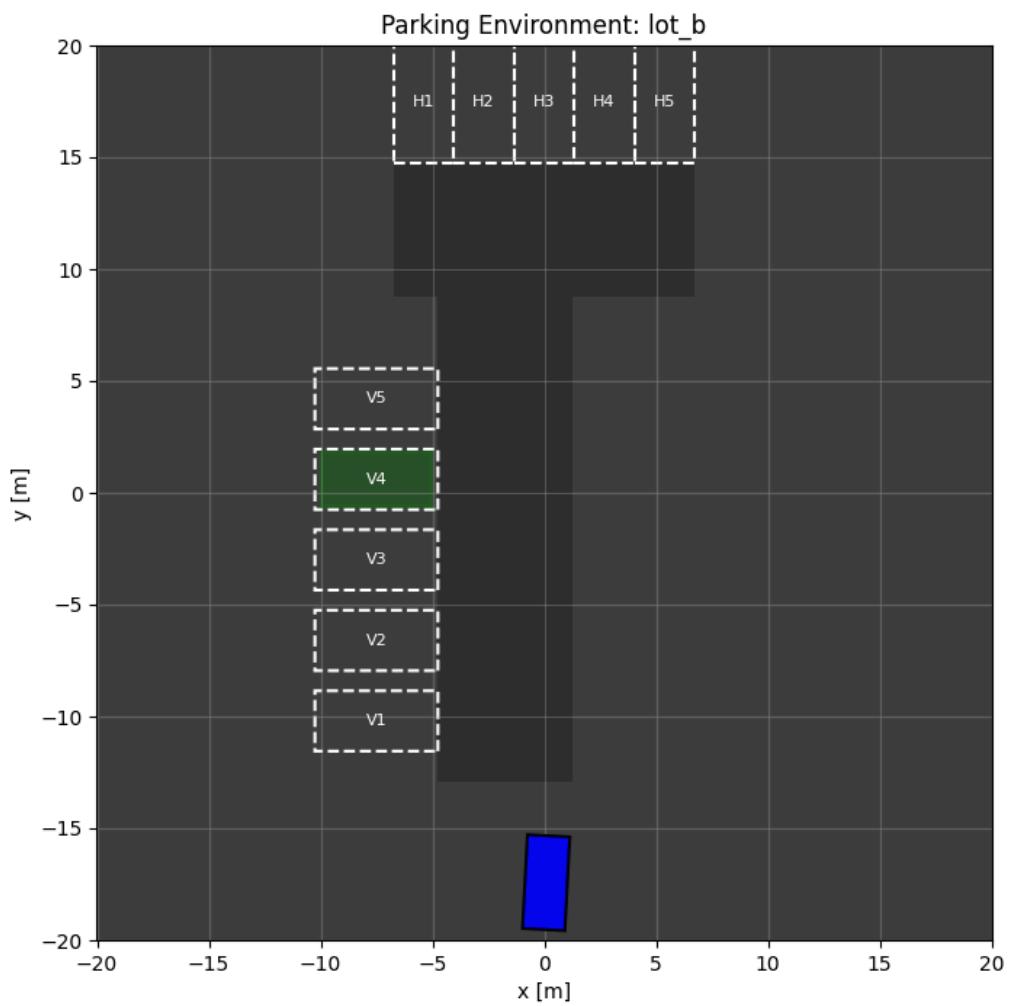


## 4.2 2D Simulation Worlds

### 2D Lot A - Matplotlib Render:



### 2D Lot B - Matplotlib Render:



## 4.3 Autonomous Controller Logs

Include **actual console outputs** from your runs (just paste as text):

## Lot A Output Example

```
(.venv) vishwaksena_dingari@vd autonomous_parking_ws % cd ~/autonomous_parking_ws
source .venv/bin/activate

python -m autonomous_parking.controller_basic --lot lot_a
Starting autonomous parking in lot_a...
Initial obs: [10.613094 -6.4020824 3.1161044 0.           12.394532   4.92
              10.          10.          ]
Step 0000 | dist=12.22 yaw_err=175.5° v=2.00 | sensors L/C/R = 3.67/10.00/10.00 | v_cmd: 10.00
Step 0020 | dist=10.99 yaw_err=123.5° v=1.90 | sensors L/C/R = 10.00/10.00/10.00 | v_cmd: 10.00
Step 0040 | dist=12.96 yaw_err=110.6° v=2.00 | sensors L/C/R = 10.00/10.00/9.81 | v_cmd: 10.00
Step 0060 | dist=15.98 yaw_err=111.1° v=2.00 | sensors L/C/R = 10.00/10.00/6.57 | v_cmd: 10.00
Step 0080 | dist=19.27 yaw_err=115.2° v=2.00 | sensors L/C/R = 10.00/10.00/5.13 | v_cmd: 10.00
Step 0100 | dist=22.63 yaw_err=120.4° v=2.00 | sensors L/C/R = 10.00/10.00/4.06 | v_cmd: 10.00
Step 0120 | dist=25.98 yaw_err=126.8° v=2.00 | sensors L/C/R = 10.00/10.00/3.25 | v_cmd: 10.00
Step 0140 | dist=29.25 yaw_err=134.4° v=2.00 | sensors L/C/R = 10.00/10.00/2.63 | v_cmd: 10.00
Step 0159 | dist=32.22 yaw_err=142.9° v=2.00 | sensors L/C/R = 10.00/10.00/2.17 | v_cmd: 10.00
Episode finished. Info: {'success': False, 'collision': True, 'dist': np.float32(32.219
Total reward: -3222.27
(.venv) vishwaksena_dingari@vd autonomous_parking_ws %
```

## Lot B Output Example

```
(.venv) vishwaksena_dingari@vd autonomous_parking_ws % cd ~/autonomous_parking_ws
source .venv/bin/activate

python -m autonomous_parking.controller_basic --lot lot_b
Starting autonomous parking in lot_b...
Initial obs: [35.08267 4.872623 -1.4471282 0. 35.41943 10.
              10. 6.07]
Step 0000 | dist=35.22 yaw_err=-79.8° v=2.00 | sensors L/C/R = 10.00/10.00/4.82 | v_cmd=0
Step 0020 | dist=32.68 yaw_err=-52.6° v=2.00 | sensors L/C/R = 10.00/10.00/10.00 | v_cm
Step 0040 | dist=29.79 yaw_err=-52.9° v=2.00 | sensors L/C/R = 10.00/10.00/10.00 | v_cm
Step 0060 | dist=27.06 yaw_err=-55.9° v=2.00 | sensors L/C/R = 10.00/10.00/10.00 | v_cm
Step 0080 | dist=24.49 yaw_err=-59.5° v=2.00 | sensors L/C/R = 10.00/10.00/10.00 | v_cm
Step 0100 | dist=22.12 yaw_err=-63.6° v=2.00 | sensors L/C/R = 10.00/10.00/10.00 | v_cm
Step 0120 | dist=19.98 yaw_err=-68.4° v=2.00 | sensors L/C/R = 10.00/10.00/10.00 | v_cm
Step 0140 | dist=18.13 yaw_err=-74.0° v=2.00 | sensors L/C/R = 10.00/10.00/10.00 | v_cm
Step 0160 | dist=16.64 yaw_err=-79.9° v=2.00 | sensors L/C/R = 10.00/10.00/8.05 | v_cmd=0
Step 0180 | dist=15.69 yaw_err=-83.7° v=2.00 | sensors L/C/R = 8.16/7.05/4.06 | v_cm=2
Step 0200 | dist=15.59 yaw_err=-79.1° v=2.00 | sensors L/C/R = 4.84/3.38/1.00 | v_cm=2
Step 0220 | dist=15.77 yaw_err=-75.3° v=0.30 | sensors L/C/R = 4.29/2.85/0.65 | v_cm=0
Step 0240 | dist=15.97 yaw_err=-73.5° v=0.30 | sensors L/C/R = 3.83/2.37/0.27 | v_cm=0
Step 0260 | dist=16.10 yaw_err=-72.6° v=0.15 | sensors L/C/R = 3.54/2.06/0.02 | v_cm=0
Step 0280 | dist=16.43 yaw_err=-87.6° v=0.15 | sensors L/C/R = 1.95/0.08/10.00 | v_cm=0
Step 0299 | dist=16.54 yaw_err=-103.1° v=0.15 | sensors L/C/R = 0.09/10.00/10.00 | v_cm=0
Episode finished. Info: {'success': False, 'collision': False, 'dist': np.float32(16.54)
Total reward: -6408.86
(.venv) vishwaksena_dingari@vd autonomous_parking_ws %
```

## 5. Issues Encountered (and fixes)

### 1. Premature episode termination

Caused by sensor < 0.2 m.

- Fixed by using soft penalties instead of hard collision.

### 2. Sensor rays detecting world boundaries too aggressively

- Fixed by lowering threshold and removing termination.

### 3. Controller getting stuck near walls

- Reduced sideways bias and tuned speed control.

#### 4. Lot A controller not reaching bays directly

- Adjusted speed and steering blending.
- Now drives into the correct area without crashing.

These fixes make the environment stable for Phase 3.

## 6. Plans for Phase 3 (Major Progress Stage)

### 6.1 Controller I Will Implement in Phase 3

For Phase 3, I will implement a **simple PPO reinforcement learning agent** using the 8-dimensional observation space from Phase 2. PPO is stable, easy to train, and well-suited for continuous control.

#### RL Setup (planned):

- **State:** [local\_x, local\_y, yaw\_err, v, dist, s\_left, s\_center, s\_right]
- **Action:** continuous [v\_cmd, steer\_cmd]
- **Goal:** reduce distance + yaw error, avoid boundaries, stop inside bay.

Phase 3 deliverables:

- PPO training loop (PyTorch)
- Rollout collection and reward logging
- Training curves (reward vs steps)
- A saved PPO checkpoint
- Evaluation runs in Lot A and Lot B

If time permits, I may also add a simple geometric controller, but PPO will be the main method for Phase 3.

## 6.2 Episode Recording + Metrics

I will add:

- success rate
- distance-to-bay curves
- yaw error per step
- reward curve
- path visualization
- trajectory plots

## 6.3 Integrate RL Training Loop

If PPO is chosen:

- vectorized environment
- replay buffer
- training logs
- policy checkpoint saving
- final evaluation video

## 6.4 Optional Extensions

If time permits:

- sync 2D + 3D visualization
- add real LiDAR plugin in Gazebo
- generate real trajectories into 3D

# 7. Conclusion

Phase 2 is completed successfully.

The autonomous parking project now has:

- fully stable 2D + 3D environments
- improved observation space
- a complete reward system
- soft sensor penalties
- safe termination conditions
- a working autonomous controller demonstration
- successful runs in both Lot A and Lot B
- clear logs and behavior explaining controller actions

The environment is now ready for **algorithm integration and training** for Phase 3.